

# Recommender System for Compass Data

24th July 2018

Ravi Raju Krishna



# Descriptive Analytics

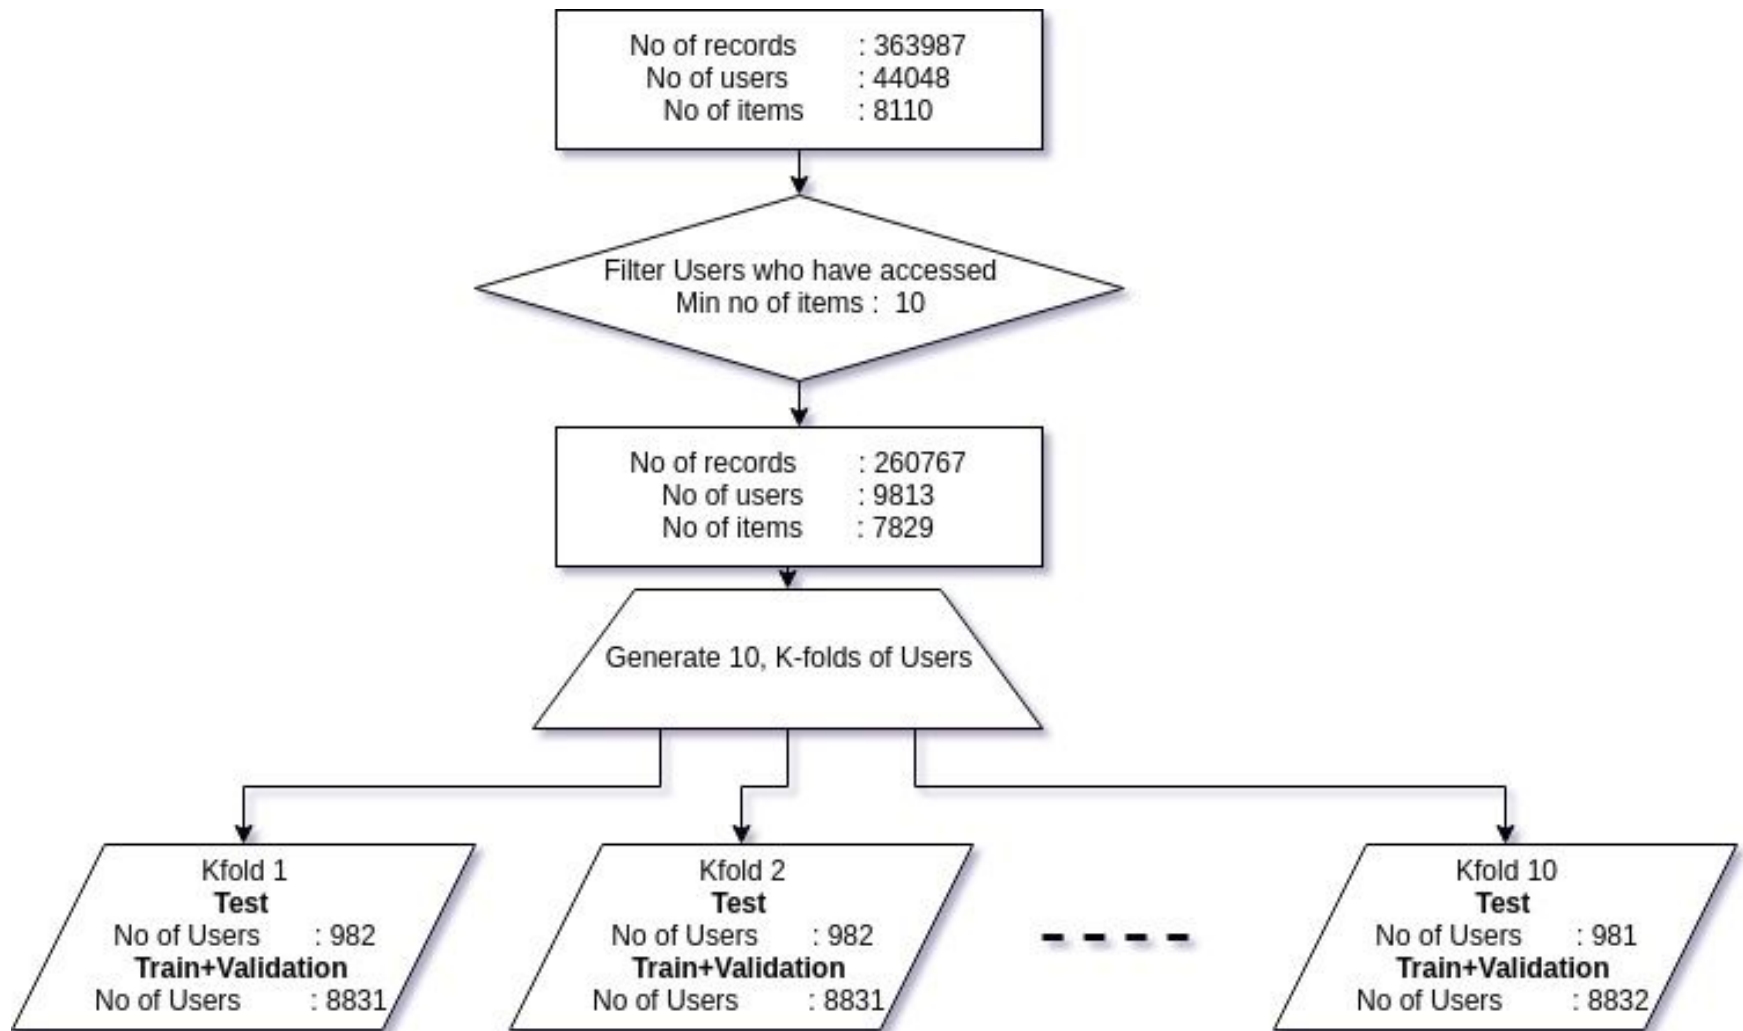
# Data Preprocessing

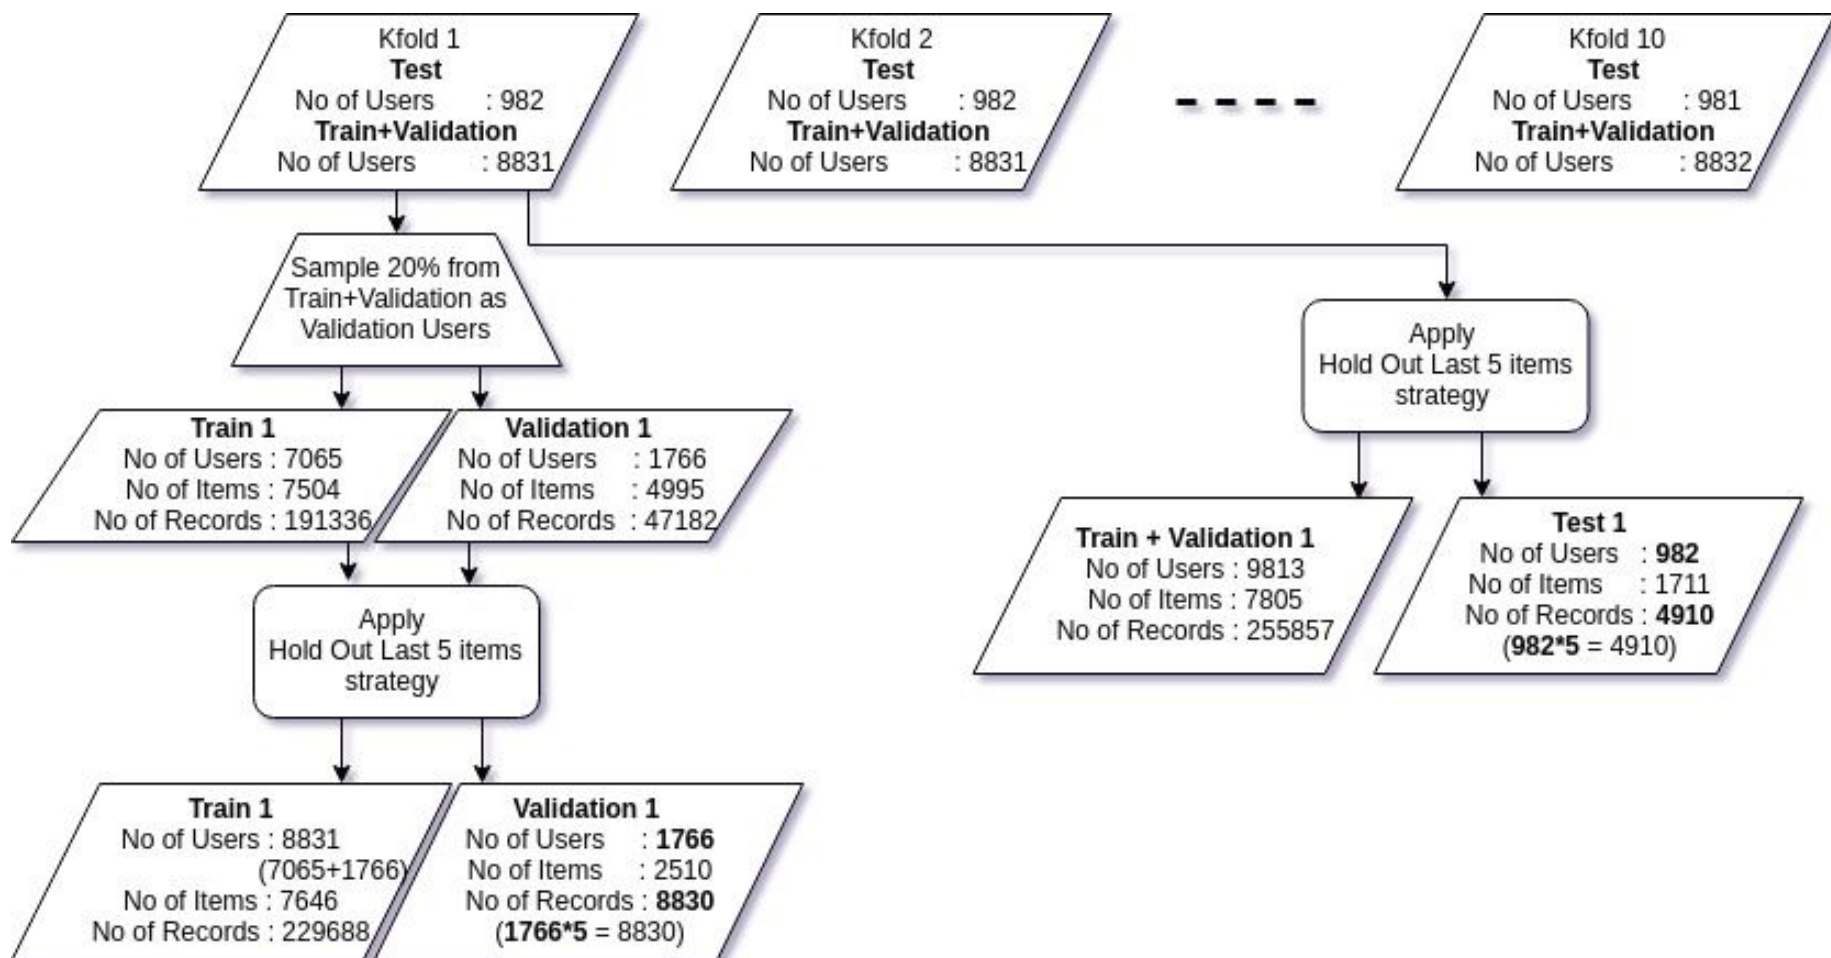
- Unique key for item : “media\_id”, Unique key for user: “learner\_id”  
Rating of each learner to media : “like\_rating”
- Total:  
No of learners : 44048    No of videos    : 8110    No of events : 424659
- Events from 2016-06-07 to 2018-04-27
- Learners provide like\_rating [1,2,3] scale
- Few (**3** events) outliers with 0 rating was filtered
- No of user,item pairs with multiple ratings : **60669**  
Identical Ratings : 55423, Distinct Ratings : 5246
- Hence only latest rating w.r.t. event\_time were chosen
- Total:  
No of learners : 44048    No of videos    : 8110    No of events : 363987  
(14% : **-60672**)

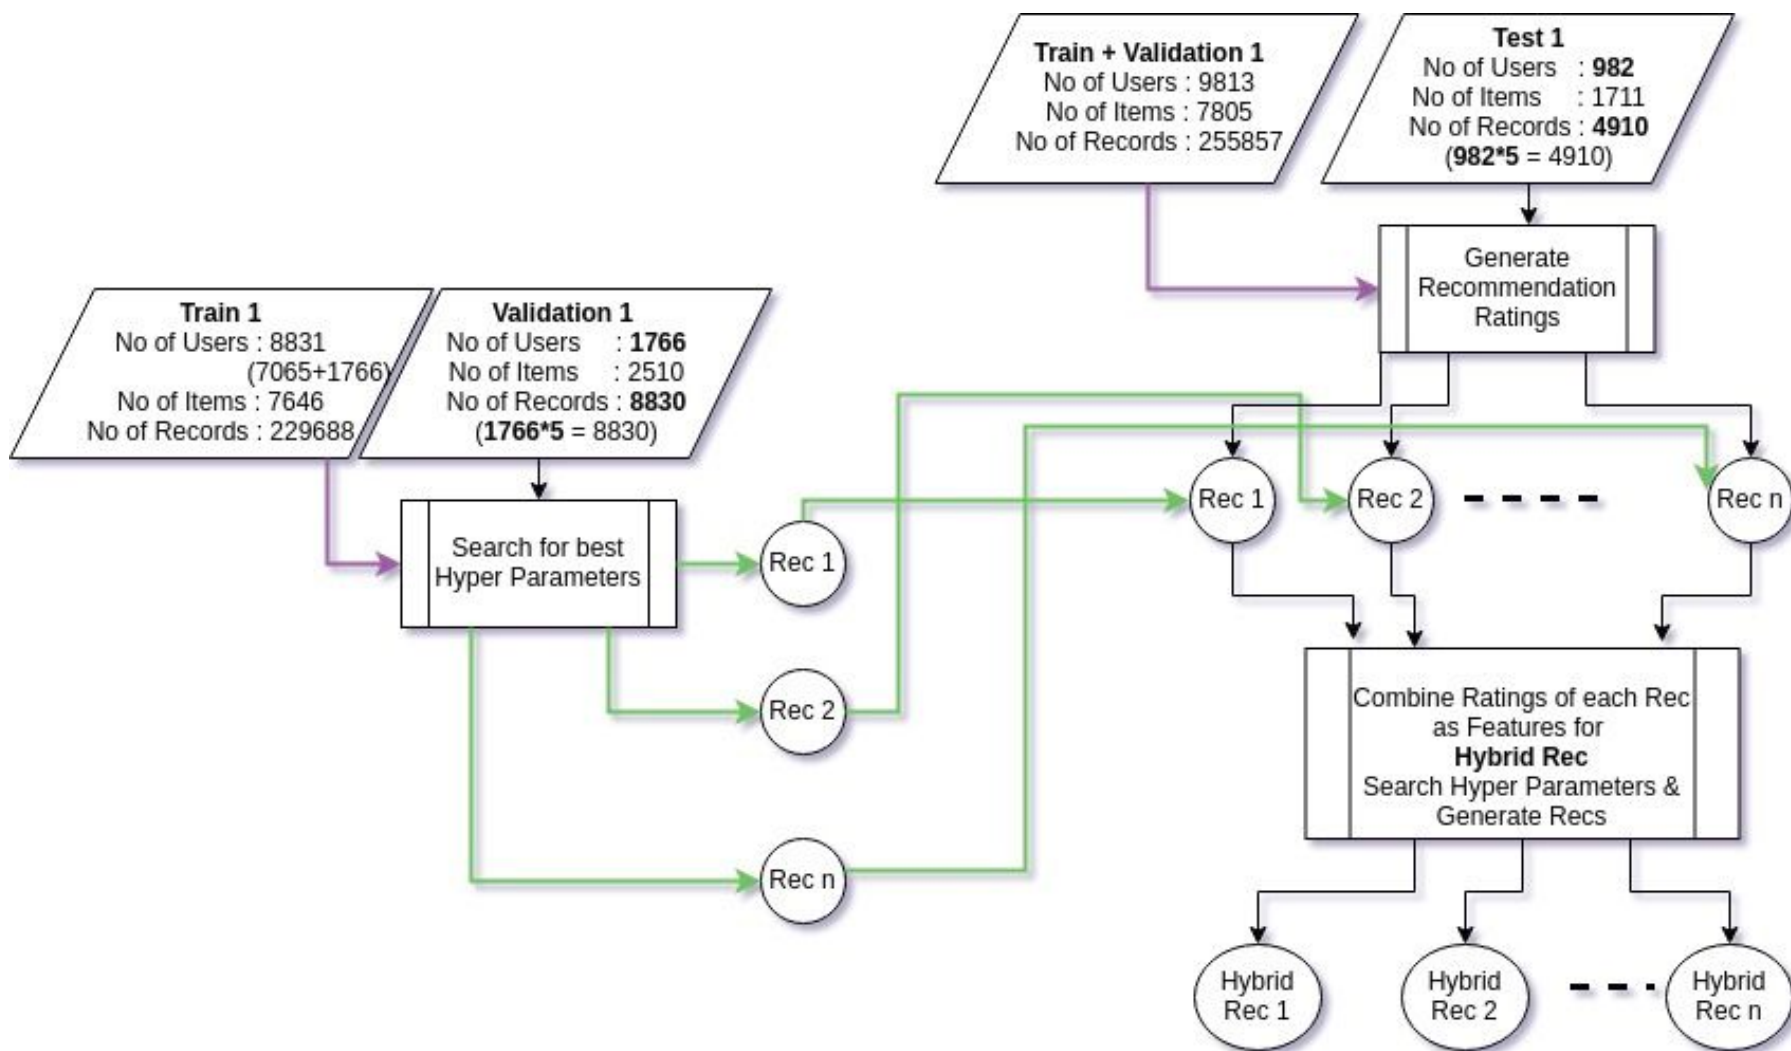
# Train, Validation and Test Data

- Total:  
No of learners : 44048    No of videos : 8110    No of events : 363987
- Filter Users who have accessed min of 10 items
- Total:  
No of learners : 9813    No of videos : 7829    No of events : 260767  
(28% : -103220)
- Generate 10 k-folds of users, each fold having ~981 users as test and ~ 8832 as train+validation users
- Randomly sample 20% of train+validation users as validation
- Hold out last 5 items of validation users as validation data, adding n-5 items to training in each fold.
- Similarly apply hold\_out\_strategy for each k-fold of test users.









# Open Source Libraries

- Surprise Lib <http://surpriselib.com/>

Simple Python **R**ecommendation **S**ystem **E**ngine.

Surprise is a python scikit library for building and analyzing rating based recommender systems.

Provide various ready-to-use prediction algorithms such as baseline algorithms, neighborhood methods, matrix factorization-based ( SVD, PMF, SVD++, NMF). Various similarity measures (cosine, MSD, pearson...) are built-in.

- ML Flow <https://mlflow.org/>

The MLflow Tracking component is an API and UI for logging parameters, code versions, metrics, and output files when running your machine learning code



# Recommendation Algorithms (1/3)

- BaselineOnly

Typical CF data exhibit large user and item effects—systematic tendencies for some users to give higher ratings than others—and for some items to receive higher ratings than others

BaselineOnly is an algorithm to predict the baseline estimate for given user and item.

$$b_{ui} = \mu + b_u + b_i.$$

The parameters  $\mathbf{b}_u$  and  $\mathbf{b}_i$  indicate the observed deviations of user  $u$  and item  $i$ , respectively, from the average.

If user/item is unknown, then the respective bias is assumed to be zero.

To estimate  $\mathbf{b}_u$  and  $\mathbf{b}_i$  the least squares problem is solved using ALS or SGD

$$\min_{b_*} \sum_{(u,i) \in \mathcal{K}} (r_{ui} - \mu - b_u - b_i)^2 + \lambda_1 \left( \sum_u b_u^2 + \sum_i b_i^2 \right)$$

# Recommendation Algorithms (2/3)

- KNNBaseline

User based methods estimate unknown ratings based on recorded ratings of like-minded users. In Item based methods, a rating is estimated using known ratings made by the same user on similar items.

Pearson correlation coefficient, measures the tendency of users to rate items  $i$  and  $j$  similarly. **Pearson-baseline correlation coefficient** between all pairs of items (or users) uses baselines for centering instead of means.

$$\text{pearson\_baseline\_sim}(i, j) = \hat{\rho}_{ij} = \frac{\sum_{u \in U_{ij}} (r_{ui} - b_{ui}) \cdot (r_{uj} - b_{uj})}{\sqrt{\sum_{u \in U_{ij}} (r_{ui} - b_{ui})^2} \cdot \sqrt{\sum_{u \in U_{ij}} (r_{uj} - b_{uj})^2}}$$

Since many ratings are unknown, some items may share only a handful of common observed raters. The empirical correlation coefficient, is based only on the common user support. Similarities based on a greater user support are more reliable.

# KNN Baseline

The **shrunk Pearson-baseline correlation coefficient** is defined as:

$$\text{pearson\_baseline\_shrunk\_sim}(i, j) = \frac{|U_{ij}| - 1}{|U_{ij}| - 1 + \text{shrinkage}} \cdot \hat{\rho}_{ij}$$

**U<sub>ij</sub>** : the set of all users that have rated both items i and j.

**shrinkage** parameter helps to avoid overfitting when only few ratings are available

Using the similarity measure, we identify the k items rated by u, which are most similar to i. This set of k neighbors is denoted by **N<sub>u</sub><sup>k</sup>(i)**

The predicted value of **r<sub>ui</sub>** is taken as a weighted average of the ratings of neighboring items, while adjusting for user and item effects through the baseline estimates

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N_u^k(j)} \text{sim}(i, j)}$$

# Recommendation Algorithms (3/3)

- Latent Factor Models

Latent factor models comprise an alternative approach to collaborative filtering with the more holistic goal to uncover latent features that explain observed ratings.

A typical model associates each user  $u$  with a user-factors vector  $\mathbf{p}_u \in \mathbf{R}^f$ , and each item  $i$  with an item-factors vector  $\mathbf{q}_i \in \mathbf{R}^f$ .

The prediction is done by taking an inner product:

$$\hat{r}_{ui} = \mu + b_u + b_i + \mathbf{q}_i^T \mathbf{p}_u$$

To estimate all the unknown, the following regularized squared error is minimized:

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|\mathbf{q}_i\|^2 + \|\mathbf{p}_u\|^2)$$

# SVD algorithm

The minimization is performed by a very straightforward stochastic gradient descent:

$$\begin{aligned}b_u &\leftarrow b_u + \gamma(e_{ui} - \lambda b_u) \\b_i &\leftarrow b_i + \gamma(e_{ui} - \lambda b_i) \\p_u &\leftarrow p_u + \gamma(e_{ui} \cdot q_i - \lambda p_u) \\q_i &\leftarrow q_i + \gamma(e_{ui} \cdot p_u - \lambda q_i)\end{aligned}$$

where  $e_{ui} = r_{ui} - \hat{r}_{ui}$

These steps are performed over all the ratings of the trainset and repeated `n_epochs` times.

Baselines are initialized to 0. User and item factors are randomly initialized according to a normal distribution



# SVD++ algorithm

The SVD++ algorithm, an extension of SVD taking into account implicit ratings.

The prediction rating is set as:

$$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left( p_u + |I_u|^{-\frac{1}{2}} \sum_{j \in I_u} y_j \right)$$

Where

$I_u$  : the set of all items rated by user  $u$

$y_j$  terms are a new set of item factors that capture implicit ratings.

An implicit rating describes the fact that a user  $u$  rated an item  $j$ , regardless of the rating value.

# Hybrid Recommendation Algorithms

Baseline\_Only, Knn\_UserBased\_Baseline, Knn\_ItemBased\_Baseline, SVD\_biased and SVDpp\_biased algorithms were used to generate recommenders for last 5 items held out for 9813 users.

These generated recommendations were used as features of the following hybrid recommenders:

- Elastic Net Linear Regression
- Gradient Boosting Regressor

# Hybrid Recommendation Algorithms (1/2)

- Elastic Net Linear Regression

Elastic Net is a linear regression model trained with L1 and L2 prior as regularizer. This combination allows for learning a sparse model where few of the weights are non-zero like Lasso, while still maintaining the regularization properties of Ridge.

The objective function to minimize is

$$\min_w \frac{1}{2n_{\text{samples}}} \|Xw - y\|_2^2 + \alpha \rho \|w\|_1 + \frac{\alpha(1 - \rho)}{2} \|w\|_2^2$$

## Hybrid Recommendation Algorithms (2/2)

- Gradient Boosting Regressor

Gradient Boosted Regression Trees (GBRT) is a generalization of boosting to arbitrary differentiable loss functions.

The number of weak learners (i.e. regression trees) is controlled by the parameter **n\_estimators**; The size of each tree can be controlled either by setting the tree depth via **max\_depth**. The **learning\_rate** is a hyper-parameter in the range (0.0, 1.0] that controls overfitting

# H<sub>2</sub>O Flow Tracking



Default

Experiment ID: 0      Artifact Location: mlruns/0

Search Runs:

metrics.rmse < 1 and params.model = "tree"

Search

Filter Params:

alpha, lr

Filter Metrics:




rmse, r2

Clear

14 matching runs

Compare Selected

Download CSV

				Parameters						Metrics	
Date	User	Source	Version	algo	best_params	no_of_experiments	param_grid	time_taken	time_taken_per_exp	best_rmse	best_rmse
 2018-07-19 07:29:42	ravi	hybrid_recommender_search.py	b08725	GradientBoostingRegressor	{'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 100}	36	{'learning_rate': [0.001, 0.01, 0.1], 'n_estimators': [10, 100, 150], 'max_depth': [1, 7, 10, 15]}	14.08 min	23.46 sec	-0.473	0.688
 2018-07-19 07:24:01	ravi	hybrid_recommender_search.py	b08725	RandomForestRegressor	{'max_depth': 7, 'n_estimators': 500}	12	{'n_estimators': [150, 250, 500], 'max_depth': [1, 7, 10, 15]}	05.69 min	28.46 sec	-0.474	0.688
 2018-07-19 07:23:10	ravi	hybrid_recommender_search.py	b08725	SGDRegressor	{'alpha': 0.0001, 'l1_ratio': 0.2, 'loss': 'squared_loss', 'penalty': 'l2', 'tol': 1e-06}	420	{'loss': ['squared_loss', 'huber'], 'penalty': ['l1', 'l2', 'elasticnet'], 'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100], 'l1_ratio':	51.03 sec	00.12 sec	-0.475	0.689

# Results

# Recommenders

Algorithm	AVG RMSE of 10-fold Validation Data	AVG RMSE of 10-fold Testing Data
<b>Knn_UserBased_Baseline_SGD_Tuned</b>	0.6909	0.6905
<b>BaselineOnly_SGD_Tuned</b>	0.6944	0.6939
<b>Knn_ItemBased_Baseline_SGD_Tuned</b>	0.6943	0.6940
Knn_UserBased_Baseline_ALS_Tuned	0.7019	0.7018
BaselineOnly_ALS_Tuned	0.7036	0.7036
Knn_ItemBased_Baseline_ALS_Tuned	0.7035	0.7036
<b>SVDpp_biased_Tuned</b>	0.7067	0.7066
<b>SVD_biased_Tuned</b>	0.7076	0.7075
SVD_unbiased_Tuned	0.7580	0.7538

# Hybrid Recommenders

<b>Meta Learner</b>	<b>AVG RMSE of 10-fold Testing Data</b>
GradientBoostingRegressor	0.6866
RandomForestRegressor	0.6869
ElasticNet	0.6880
LinearRegression	0.6880

# Comparison with existing recommender

- In order to compare with existing recommender in production, Anti-Test set(User/Item Pairs not in input event data) was generated.
- Individual recommender algorithms and hybrid recommenders were used to predict ratings for user/item pairs in anti-test set and top 50 recommendations were generated for each user.

Recommender	RMSE
hybrid_Gradient Boosting Regressor	0.6291
hybrid_Elastic Net Regressor	0.6300
Knn_UserBased_Baseline_SGD	0.6378
BaselineOnly_SGD	0.6385
Knn_ItemBased_Baseline_SGD	0.6437
<b>current</b>	<b>0.6752</b>
SVD ++	0.6875
SVD	0.6877



# Enhancements

1. Current recommenders used only latest rating, leverage multiple rating events, analyse whether time duration spent on video impacts change in rating.
2. Use other metadata of events like difficult\_rating, playlist/channel information
3. Aggregation of events data to analyse no of repeated access in addition to time spent, to associate confidence value for rating
4. Incorporate Temporal Dynamics effect on user, item biases and user preferences to model rating as function of time
5. Average time between events to get heuristics of engagement.



*Personalized, Playful, Lifelong Learning*